

Plant Disease Classifier

Group 23

Asım Adil Can - 22103729

Erdem Pülat - 22103566

Gün Taştan - 22101850

Kerem Er - 22102369

Begüm Kunaç - 22103838



Project Overview

Significance:

Early detection and accurate classification of plant diseases are essential for timely intervention and effective disease management. By developing a robust classification system that identifies plant health conditions, our project aims to:

- Improve plant health monitoring.
- Enable early intervention to minimize crop loss.
- Enhance agricultural efficiency through automation

Objective:

To develop a machine learning-based classification system capable of accurately identifying three distinct plant conditions:

- **Healthy Plants:** No visible disease symptoms.
- **Powdery Mildew:** Characterized by white fungal growth on leaves.
- **Rust Disease:** Identified by yellow, rust-like discoloration on leaves.

Our goal is to compare multiple machine learning and deep learning approaches to determine the most effective method for this classification task.

Dataset

Class Distribution:

- **Healthy:** Plants showing no visible disease symptoms.
- **Powdery:** Plants affected by white fungal growth.
- **Rusty:** Plants displaying yellow to orange discoloration on its surface.

Dataset Structure:

- **Training Set:** 1322 images (458 healthy, 430 rusty, 434 powdery).
- **Testing Set:** 150 images (50 healthy, 50 rusty, 50 powdery).
- **Validation Set:** 60 images (20 healthy, 20 rusty, 20 powdery).

Dataset Properties

Image Characteristics:

- **Content:** Healthy tissue, white fungal growth (powdery mildew), and yellow discoloration (rust).
- **Image Sizes:** Image sizes are not fixed, resolution of images are in good quality.

Dataset Link: [Plant Disease Classification Dataset](#)

Sample Images from Dataset



Healthy Class



Powdery Class



Rusty Class

Dataset Preprocessing

Preprocessing Pipeline:

1. Image Standardization:

- Resized all images to **128x128 pixels** for consistent input dimensions and reduced computational overhead.
- Used PyTorch's **torchvision.transforms** for resizing.

2. Normalization:

- Images normalized with mean and standard deviation of 0.5 for each RGB channel.
- Transformed pixel values to the range **[-0.5, 0.5]** to achieve faster convergence and better numerical stability.

What did we use?



```
graph TD; A[What did we use?] --> B[Libraries we used:]; A --> C[Additional tools:]; A --> D[Models we trained:];
```

Libraries we used:

- NumPy
- OpenCV
- Matplotlib
- Scikit-learn
- pandas
- Pytorch

Additional tools:

- Github
- Google Colab
- Kaggle

Models we trained:

- SVM
- Random Forest
- Adaboost Classifier
- Transfer Learning
- CNN

Support Vector Machine

Preprocessing Steps for SVM:

Feature Extraction:

- **HOG Features:** Captured texture and shape patterns from grayscale images.
- **Color Histograms:** Represented color distribution in the RGB channels.

Feature Concatenation: Combined HOG and color histogram features for a richer representation.

Normalization: Features were scaled using StandardScaler for better model performance.





Figure 1: SVM illustration


Sample Image - Healthy



HOG Features - Healthy



Color Histogram - Healthy




Sample Image - Powdery



HOG Features - Powdery



Color Histogram - Powdery




Sample Image - Rust



HOG Features - Rust



Color Histogram - Rust





SVM - Experimental Setup and Parameters

- **Grid Search:**
 - Performed hyperparameter optimization using 5-fold cross-validation.
- **Evaluation Criteria:**
 - Weighted F1-score.
- **Parameter Grid:**
 - **Kernel:** Radial Basis Function, Linear, Polynomial
 - **C (Regularization parameter):** 0.01, 0.1, 1, 10, 100
 - **Degree:** 2, 3, 4
 - **Gamma:** 0.001, 0.01, 0.1, 1, scale

SVM - Results

- **Best Parameters:**
 - **Kernel:** Radial Basis Function
 - **C (Regularization parameter)** 10
 - **Gamma:** scale
- **Validation Results:**
 - Accuracy: 0.7833
 - Weighted F1-Score: 0.7797
- **Test Results:**
 - Accuracy: 0.76
 - Weighted F1-Score: 0.7554



Random Forest

Preprocessing Steps for RF:

1. Feature Extraction:

- **HOG Features:** Captured texture and shape patterns from grayscale images.
- **Color Histograms:** Represented color distribution in the RGB channels.

2. Feature Concatenation:

Combined HOG and color histogram features for a richer representation.

3. Normalization:

Features were scaled using StandardScaler for better model performance.

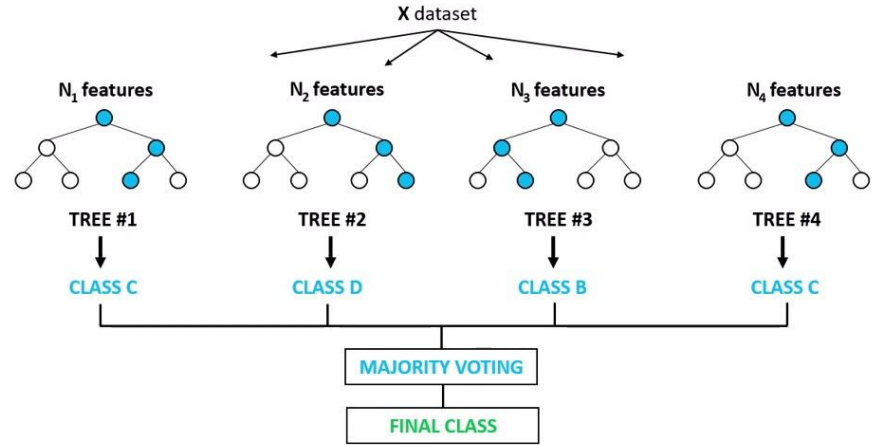




Figure 3: Random Forest Illustration

Random Forest - Experimental Setup and Parameters

- **Grid Search:**
 - Performed hyperparameter optimization using 5-fold cross-validation.
- **Evaluation Criterion:**
 - Weighted F1-score.
- **Parameter Grid:**
 - n_estimators: 100, 200, 300
 - max_depth: 10, 20, None
 - min_samples_split: 2, 5, 10
 - min_samples_leaf: 1, 2, 4
 - max_features: sqrt, log2, None
 - criterion: gini, entropy

RF Results

- **Best Parameters:**
 - criterion: gini
 - max_depth: None
 - max_features: sqrt
 - min_samples_leaf: 1
 - min_samples_split: 2
 - n_estimators: 200
- **Validation Results:**
 - Accuracy: 0.9667
 - Weighted F1-Score: 0.9666
- **Test Results:**
 - Accuracy: 0.9200
 - Weighted F1-Score: 0.9209

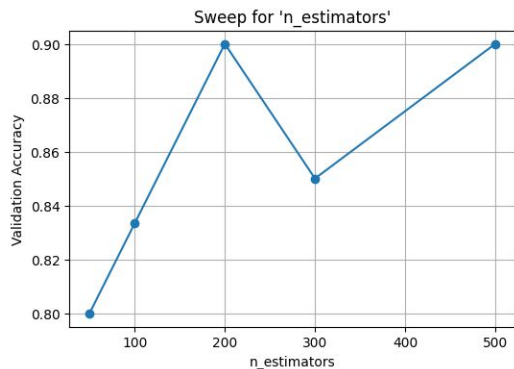


Adaboost & The Core Mechanism

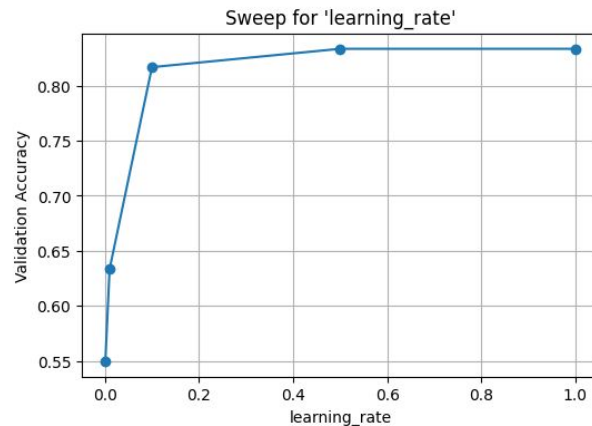
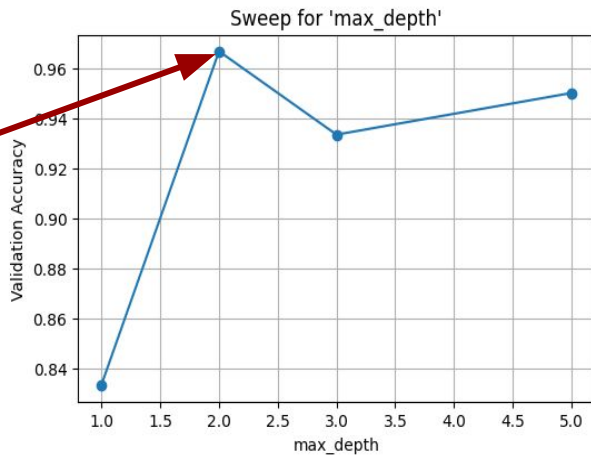
- AdaBoost stands for *Adaptive Boosting*.
- Initialize equal weights for all training examples.
- Train a weak learner; measure error on weighted samples.
- Increase weights for misclassified samples; decrease for correct ones.
- Repeat for multiple rounds and sum each learner's contributions.
- Sequentially focuses on misclassified examples to reduce errors over rounds.
- Our weak learner: Decision Trees, with depth as a hyperparameter.
- Feature Extraction: Same as in SVM, combination of grayscale edge-features and color histograms

Adaboost Hyperparameters & Tuning

Base : # of estimators : 100, Learning Rate: 1.0, Max Depth: 1




The Winner!



Adaboost Results

Test Accuracy : 86% (Tested with the winner of the hyperparameters)



Transfer learning Model

What is transfer learning?

Which Model We Used?

ResNet-18




Figure 4: Resnet 18 structure

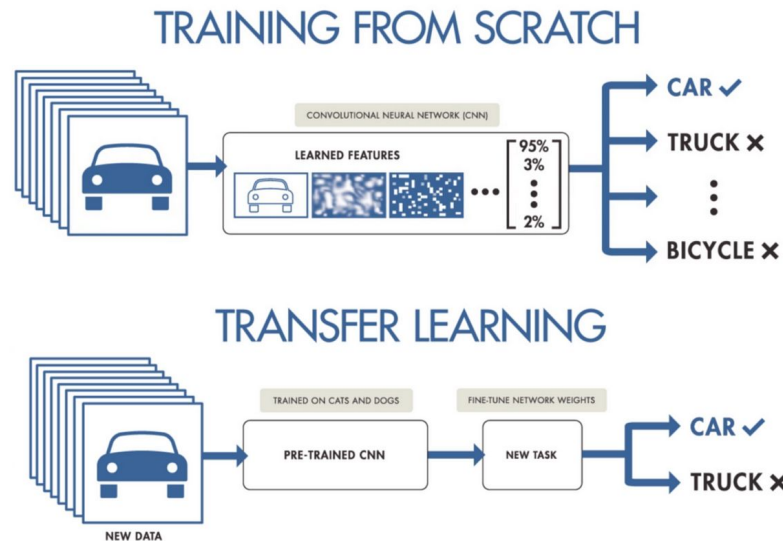


Figure 3: Transfer learning illustration


How does Transfer Learning work?

Select a Pre-trained Model

Freeze the Base Layers

Change the final output layer

Fine-Tuning



Hyperparameter tuning


	Training Loss	Validation Loss	Validation Accuracy (%)	Test Accuracy (%)
Batch Size 16	0.2	0.265	90.33	91.0
Batch Size 32	0.176	0.255	91.67	92.0
Batch Size 64	0.16	0.273	89.0	89.67

Learning rate fixed at 0.0005

Learning Rate	Epoch	Validation Loss	Validation Accuracy (%)	Test Accuracy (%)
0.0001	20.0	0.32	85.0	84.0
0.0005	20.0	0.255	91.67	92.0
0.001	20.0	0.345	88.33	86.0
0.005	20.0	0.45	83.0	81.5

Batch size fixed at 32

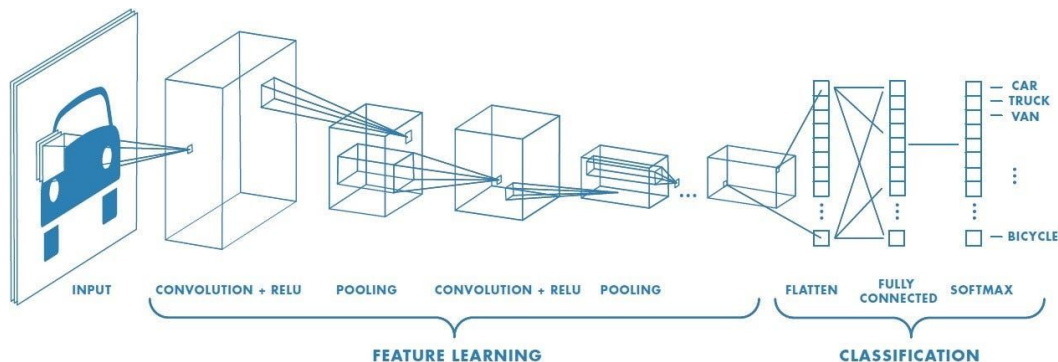
Results on transfer learning



	Precision	Recall	F1-Score	Support
Healthy	0.903846	0.94	0.921569	50.0
Powdery Mildew	0.9375	0.9	0.918367	50.0
Rust	0.92	0.92	0.92	50.0
accuracy	0.92	0.92	0.92	0.92
macro avg	0.920449	0.92	0.919979	150.0
weighted avg	0.920449	0.92	0.919979	150.0

Test Loss: 0.2942 Test Accuracy: 92.00%

CNN STRUCTURE



First Convolution Kernel

Size: $(12 \times 5 \times 5)$

Second Convolution Kernel

Size: $(18 \times 3 \times 3)$

Third Convolution Kernel

Size: $(24 \times 3 \times 3)$

Activation Function: ReLU

Pooling: Max Pooling (2×2)

Hyperparameter Tuning

Fixed Parameters During Hyperparameter Tuning:

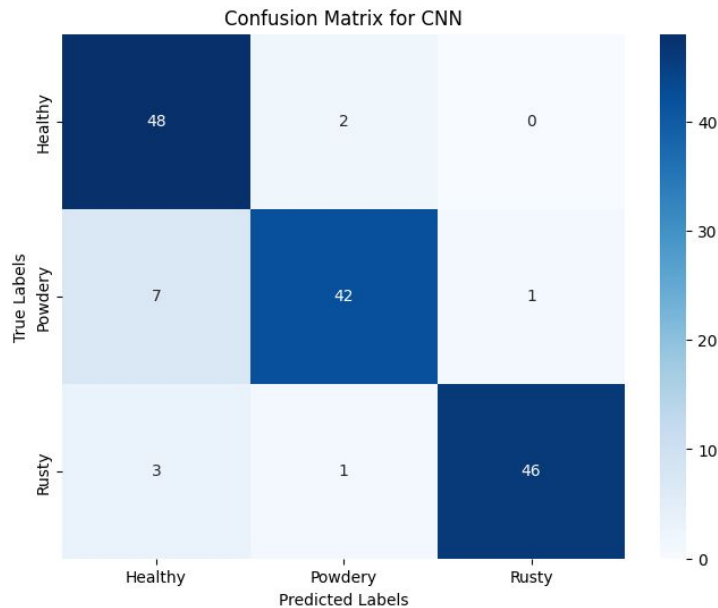
- Stochastic Gradient Descent Optimizer
- Momentum Coefficient = 0.9
- Epoch Number = 20
- Learning Rate = 0.1 in tuning of Batch Size
- Batch Size = 16 in tuning of Learning Rate

Batch Size	Training Loss	Validation Loss	Validation Accuracy	Test Accuracy
16	0.2894	0.4591	92.16%	91.44%
32	0.3071	0.4998	83.12%	83.16%
64	0.3123	0.5027	81.67%	80.44%

Learning Rate	Training Loss	Validation Loss	Validation Accuracy	Test Accuracy
0.005	0.3707	0.5218	81.94%	78.29%
0.01	0.2894	0.4593	92.00%	91.67%
0.0125	0.3092	0.4871	88.33%	87.45%

Optimal Batch Size and Learning Rate Values = (16, 0.01)


CNN RESULTS



Class	Precision	Recall	F1-Score	Support
Healthy	0.83	0.96	0.89	50
Powdery	0.93	0.82	0.88	50
Rusty	0.98	0.92	0.95	50
Macro Avg	0.91	0.91	0.91	150
Weighted Avg	0.91	0.91	0.91	150

Test Accuracy: 91.67 %

BEST Performing Model



Challenges

Limited Dataset Size



Our dataset was relatively small to other datasets, that's why we used more models to evaluate the performance

References

- [1] NeuralNine, “Image Classification CNN in PyTorch,” YouTube, <https://www.youtube.com/watch?v=CtzfbUwrYGI&t=1353s> (accessed Nov. 17, 2024).
- [2] Jorgecardete, “Convolutional Neural Networks: A comprehensive guide,” Medium, <https://medium.com/thedeephub/convolutional-neural-networks-a-comprehensive-guide-5cc0b5eae175> (accessed Nov. 17, 2024).
- [3] T. N. Fatyanosa and M. Aritsugi, “Effects of the number of hyperparameters on the performance of Ga-CNN,” *2020 IEEE/ACM International Conference on Big Data Computing, Applications and Technologies (BDCAT)*, pp. 144–153, Dec. 2020. doi:10.1109/bdcat50828.2020.00016
- [4] Ciresan, Dan & Meier, Ueli & Masci, Jonathan & Gambardella, Luca Maria & Schmidhuber, Jürgen. (2011). Flexible, High Performance Convolutional Neural Networks for Image Classification. *International Joint Conference on Artificial Intelligence IJCAI-2011*. 1237-1242. 10.5591/978-1-57735-516-8/IJCAI11-210.
- [5] “Image classification using machine learning: Support Vector Machine (SVM),” Medium, Oct. 2020. Available: <https://medium.com/analytics-vidhya/image-classification-using-machine-learning-support-vector-machine-svm-dc7a0ec92e01>. Accessed: Dec. 21, 2024.
- [6] “Image classification using sklearn random forest,” Kaggle, Oct. 2020. Available: <https://www.kaggle.com/code/kkhandekar/image-classification-using-sklearn-randomforest>. Accessed: Dec. 21, 2024.